



Advanced Operating Systems **Multi-threading in C++**

Giuseppe Massari, Federico Terraneo

giuseppe.massari@polimi.it, federico.terraneo@polimi.it

Introduction

- Multi-tasking vs multi-threading

C++11 threading support

- Thread creation
- Synchronization
- Mutex usage and pitfalls
- Condition variable

Design patterns


- Producer/Consumer
- Active Object
- Reactor
- ThreadPool

Multi-tasking vs multi-threading

- Multi-tasking operating systems allow to run more “processes” concurrently

A single application can spawn multiple processes


OS assigns a separate address space to each process
(a process cannot directly access the address space of another process)



- Multi-threading allows a single process to perform multiple tasks concurrently (in a shared address space)

Fast and easy sharing of data structures among tasks

- Both were introduced to improve performance and responsiveness of computing systems
- This class is focused on multi-threading*



Why multi-threading?

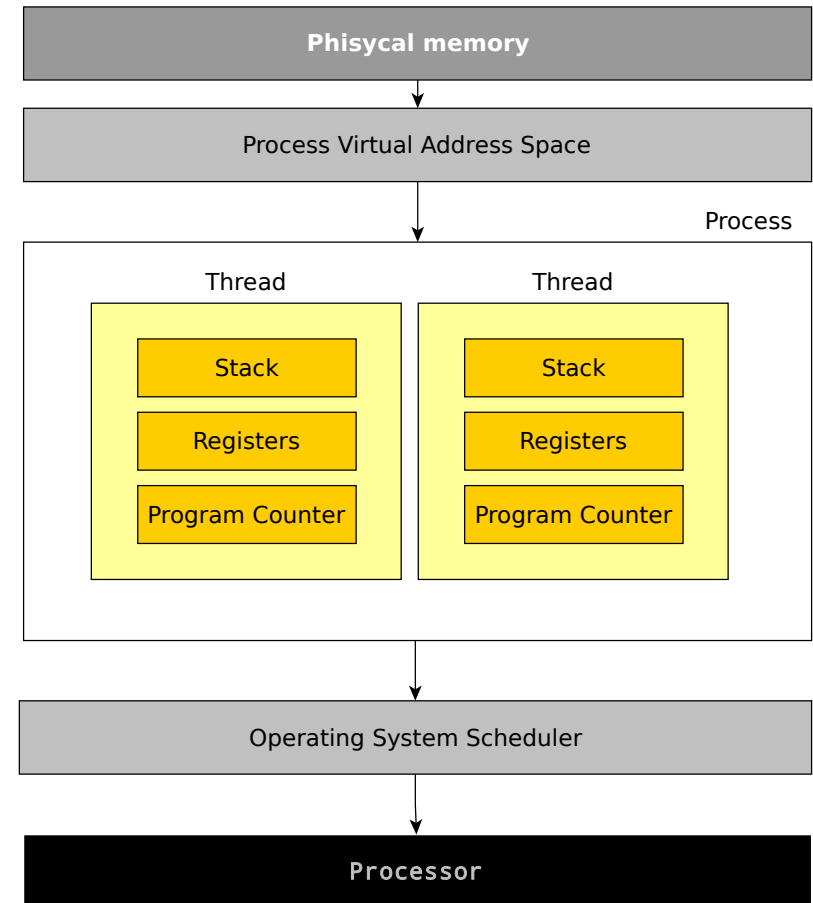
- Allowing an applications to perform multiple tasks concurrently
For example running a responsive user interface, handling network communications and running heavy batch computations at the same time
- Parallelizing algorithms in order to exploit multi-core CPU
This became popular with the widespread adoption of multi-core architectures

Multi-threading support

- HW side: we are currently in the multi-core (and many-core) era
Additional performance of computing architectures delivered through an increasing number of cores
- SW side: Growing importance in the computing landscape
Programming languages are adding native support for multi-threading
Example: C++ starting from C++11 standard version

Thread

- A thread is defined as a “lightweight task”
- Each thread has a separate stack and context (registers and program counter value)
- Depending on the implementation the OS or the language runtime are responsible of the thread-to-core scheduling



Class `std::thread`

- Constructor starts a thread executing the given function

```
#include <iostream>
#include <thread>
using namespace std;
using namespace std::chrono;

void myThread() {
    for(;;) {
        cout<<"world "<<endl;
        this_thread::sleep_for(milliseconds(500));
    }
}

int main() {
    thread t(myThread);
    for(;;) {
        cout<<"hello"<<endl;
        this_thread::sleep_for(milliseconds(500));
    }
}
```

Class `std::thread`

```
$ g++ main.cpp -o test -std=c++11 -pthread
```

- There is no guarantee about the threads execution order

```
$ ./test
helloworld

helloworld

hello
world
helloworld

hello
world
helloworld
```

Class `std::thread`

- Thread constructor can take additional arguments that are passed to the thread function
- `join()` member function waits for the thread to complete

```
#include <iostream>
#include <thread>
using namespace std;

void myThread(int i, const string & s) {
    cout<<"Called with "<<i<<" and "<<s<<endl;
}

int main() {
    thread t(myThread, 2, "test");
    t.join();
}
```


Synchronization

- What is the output of the following code?

```
#include <iostream>
#include <thread>
using namespace std;
static int sharedVariable=0;

void myThread() {
    for(int i=0;i<1000000;i++) sharedVariable++;
}

int main() {
    thread t(myThread);
    for(int i=0;i<1000000;i++) sharedVariable--;
    t.join();
    cout<<"sharedVariable="<<sharedVariable<<endl;
}
```

Synchronization

```
$ ./test
sharedVariable=-313096
$ ./test
sharedVariable=-995577
$ ./test
sharedVariable=117047
$ ./test
sharedVariable=116940
$ ./test
sharedVariable=-647018
```

- Accessing the same variables from different threads can lead to *race conditions*
- The two threads are concurrently changing the value of the same variable

Synchronization

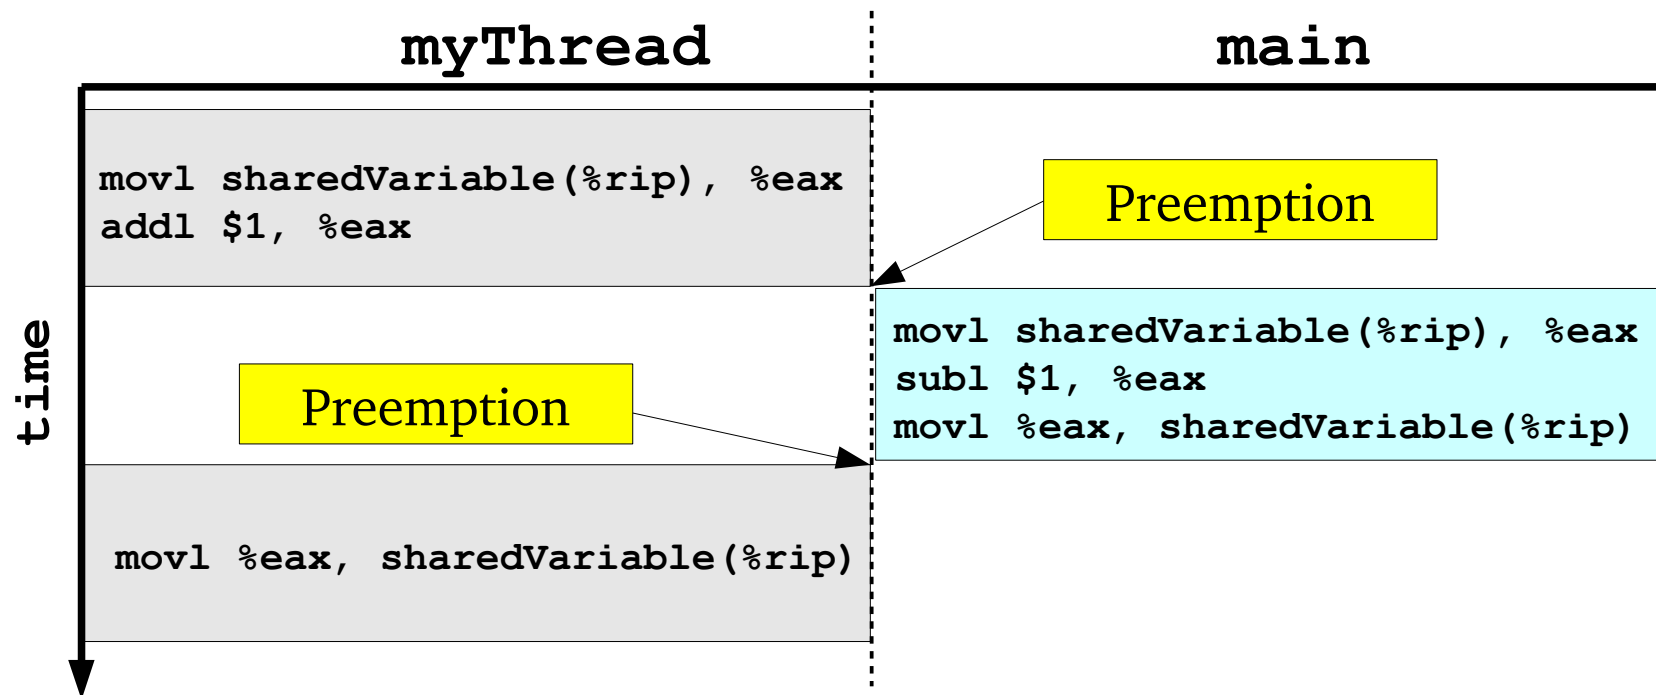
- What does happen under the hood?

```
//sharedVariable++  
movl    sharedVariable(%rip), %eax  
addl    $1, %eax  
movl    %eax, sharedVariable(%rip)  
  
//sharedVariable--  
movl    sharedVariable(%rip), %eax  
subl    $1, %eax  
movl    %eax, sharedVariable(%rip)
```

- Increment (++) and decrement (--) are not *atomic* operations
- The operating system can preempt a thread between any instruction

Synchronization

- What does happen under the hood?



- myThread** has been preempted before the result of the increment operation has been written back (**sharedVariable** update)
- This leads to incorrect and unpredictable behaviours

Synchronization

- A *critical section* is a sequence of operations accessing a shared data structure that must be performed atomically to preserve the program correctness
- In the example we faced a race condition, since the two threads entered a critical section in parallel
- To prevent race conditions we need to limit concurrent execution whenever we enter a critical section

Solution 1: *Disable preemption*

- Dangerous – it may lock the operating system
- Too restrictive – it is safe to preempt to a thread that does not modify the same data structure
- Does not work on multi-core processors

Solution 2: *Mutual exclusion*

- Before entering a critical section a thread checks if it is “free”
 - If it is, enter the critical section
 - If not, it blocks
- When exiting a critical section a thread checks if there are other blocked threads
 - If yes, one of them is selected and woken up

Class `std::mutex`

- It has two member functions

`lock()` to call before entering a critical section

`unlock()` to call after exiting a critical section

```
#include <iostream>
#include <thread>
#include <mutex>
using namespace std;
static int sharedVariable=0;
mutex myMutex;

void myThread() {
    for(int i=0;i<1000000;i++) {
        myMutex.lock();
        sharedVariable++;
        myMutex.unlock();
    }
}
```

```
int main() {
    thread t(myThread);
    for(int i=0;i<1000000;i++) {
        myMutex.lock();
        sharedVariable--;
        myMutex.unlock();
    }
    t.join();
    cout<<"sharedVariable="
        <<sharedVariable<<endl;
}
```

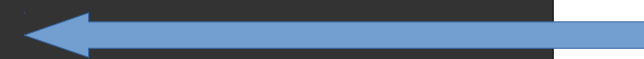
Deadlock

- Improper use of mutex may lead to *deadlock*, according to which program execution get stuck

Deadlocks may occur due to several causes

- Cause 1: *Forgetting to unlock a mutex*

```
...  
mutex myMutex;  
int sharedVariable;  
  
void myFunction(int value) {  
    myMutex.lock();  
    if(value<0) {  
        cout<<"Error"<<endl;  
        return;  
    }  
    SharedVariable += value;  
    myMutex.unlock();  
}
```


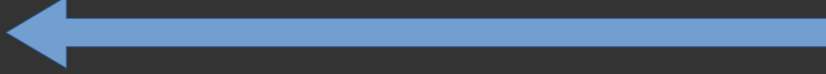


A function returns
without unlocking the
previously locked
mutex

Next function call will
result in a deadlock

Deadlock

- Solution: C++11 provides *scoped lock* that automatically unlocks mutex, regardless of how the scope is exited

```
...  
mutex myMutex;  
int sharedVariable;  
  
void myFunction(int value) {  
    {  
        lock_guard<mutex> lck(myMutex);  
        if(value<0)  
        {  
            cout<<"Error"<<endl;  
            return;   
        }  
        SharedVariable += value;  
    }   
}
```

myMutex unlocked here

OR

myMutex unlocked here

Deadlock

- Cause 2: *Nested function calls locking the same mutex*

```
...  
mutex myMutex;  
int sharedVariable;
```

```
void func2 ()
```

```
{
```

```
    lock_guard<mutex> lck (myMutex) ;  
    doSomething2 () ;
```

```
}
```

```
void func1 ()
```

```
{
```

```
    lock_guard<mutex> lck (myMutex) ;  
    doSomething1 () ;  
    func2 () ;
```

```
}
```

Deadlock if we're called
by **func1 ()**

Called with the mutex
locked

Deadlock

- Solution: *Recursive mutex* permit multiple locks by the *same thread*

```
...
recursive_mutex myMutex;
int sharedVariable;

void func2() {
    lock_guard<recursive_mutex> lck (myMutex);
    doSomething2();
}

void func1() {
    lock_guard<recursive_mutex> lck (myMutex);
    doSomething1();
    func2();
}
```

- However recursive mutex are more expensive than mutex
Use them only when needed

Deadlock

- Cause 3: *Order of locking of multiple mutexes*

```
...
mutex myMutex1;
mutex myMutex2;

void func2() {
    lock_guard<mutex> lck1(myMutex1);
    lock_guard<mutex> lck2(myMutex2);
    doSomething2();
}

void func1() {
    lock_guard<mutex> lck1(myMutex2);
    lock_guard<mutex> lck2(myMutex1);
    doSomething1();
}
```

- Thread1 calls **func1()**, locks **myMutex2** and blocks on **myMutex1**
- Thread2 calls **func2()**, locks **myMutex1** and blocks on **myMutex2**

Deadlock

- Solution: C++11 **lock** function takes care of the correct locking order

```
mutex myMutex1;  
mutex myMutex2;  
  
void func2() {  
    lock(myMutex1, myMutex2);  
    doSomething2();  
    myMutex1.unlock();  
    myMutex2.unlock();  
}  
  
void func1() {  
    lock(myMutex2, myMutex1);  
    doSomething1();  
    myMutex2.unlock();  
    myMutex1.unlock();  
}
```

- Any number of mutexes can be passed to **lock** and in any order
Use of **lock** is more expensive of **lock_guard**

Deadlocks and race conditions

- Are faults that occur due to some “unexpected” order of execution of the threads

Correct programs should work regardless of the execution order

- The order that triggers the fault can be **extremely** uncommon

Running the same program million of times may still not trigger the fault

- Are thus hard to debug

It is difficult to reproduce the bug

- Testing is almost useless for checking such errors

Good design is mandatory

Loosing concurrency

- Leaving a mutex locked for a long time reduces the concurrency in the program

```
...  
  
mutex myMutex;  
int sharedVariable=0;  
  
void myFunction()  
{  
    lock_guard<mutex> lck(myMutex);  
    sharedVariable++;  
    this_thread::sleep_for(milliseconds(500));  
}
```

Loosing concurrency

- Solution: Keep critical sections as short as possible

Leave unnecessary operations out of the critical section

```
...
mutex myMutex;
int sharedVariable=0;

void myFunction()
{
    {
        lock_guard<mutex> lck(myMutex);
        sharedVariable++;
    }
    this_thread::sleep_for(milliseconds(500));
}
```


Condition variable

- In many multi-threaded programs we may have dependencies among threads
- A “dependency” can come from the fact that the thread must wait for another thread to complete its current operation
- In such a case we need a mechanism to explicitly block a thread

class `std::condition_variable`

- Three member function
 - `wait(unique_lock<mutex> &)` - block the thread until another thread wakes it up. The mutex is unlocked for the duration of the `wait(...)`
 - `notify_one()` - wake up one of the waiting threads
 - `notify_all()` - wake up all the waiting threads
 - If no thread is waiting do nothing

class `std::condition_variable`

- In the example, `myThread` is waiting for the `main` to complete the read from standard input

```
#include <iostream>
#include <thread>
#include <mutex>
#include <condition_variable>
using namespace std;
string shared;
mutex myMutex;
condition_variable myCv;

void myThread() {
    unique_lock<mutex> lck (myMutex);
    if (shared.empty()) myCv.wait(lck);
    cout<<shared<<endl;
}
```

```
int main() {
    thread t(myThread);
    string s;
    cin>>s; // read from stdin
    {
        unique_lock<mutex>
            lck (myMutex);
        shared=s;
        myCv.notify_one();
    }
    t.join();
}
```

Introduction

- Multi-tasking vs multi-threading


C++11 threading support

- Thread creation
- Synchronization
- Mutex usage and pitfalls
- Condition variable

Design patterns

- Producer/Consumer
- Active Object
- Reactor
- ThreadPool

Producer/Consumer



- A thread (consumer) needs data from another thread (producer)
- To decouple the operations of the two threads we put a queue between them, to buffer data if the producer is faster than consumer
- The access to the queue needs to be synchronized

Not only using a mutex but the consumer needs to wait if the queue is empty

Optionally, the producer may block if the queue is full

Producer/Consumer

▪ `synchronized_queue.h` (1/2)

```
#ifndef SYNC_QUEUE_H_
#define SYNC_QUEUE_H_
#include <list>
#include <mutex>
#include <condition_variable>
template<typename T>
class SynchronizedQueue {
public:
    SynchronizedQueue() {}
    void put(const T & data);
    T get();
private:
    SynchronizedQueue(const SynchronizedQueue &);
    SynchronizedQueue & operator=(const SynchronizedQueue &);
    std::list<T> queue;
    std::mutex myMutex;
    std::conditionVariable myCv;
};
...
```

Producer/Consumer

▪ `synchronized_queue.h` (2/2)

```
template<typename T>
void SynchronizedQueue<T>::put(const T& data)
{
    std::unique_lock<std::mutex> lck(myMutex);
    queue.push_back(data);
    myCv.notify_one();
}

template<typename T>
T SynchronizedQueue<T>::get()
{
    std::unique_lock<std::mutex> lck(myMutex);
    while(queue.empty())
        myCv.wait(lck);
    T result=queue.front();
    queue.pop_front();
    return result;
}
#endif // SYNC_QUEUE_H_
```

Producer/Consumer

▪ main.cpp

```
#include "synchronized_queue.h"
#include <iostream>
#include <thread>
using namespace std;
using namespace std::chrono;

SynchronizedQueue<int> queue;

void myThread() {
    for(;;) cout<<queue.get()<<endl;
}

int main() {
    thread t(myThread);
    for(int i=0;;i++) {
        queue.put(i);
        this_thread::sleep_for(seconds(1));
    }
}
```

Producer/Consumer

- What if we do not use the condition variable?
- `synchronized_queue.h` (1/2)

```
#ifndef SYNC_QUEUE_H_
#define SYNC_QUEUE_H_
#include <list>
#include <mutex>
template<typename T>
class SynchronizedQueue {
public:
    SynchronizedQueue() {}
    void put(const T & data);
    T get();
private:
    SynchronizedQueue(const SynchronizedQueue &);
    SynchronizedQueue & operator=(const SynchronizedQueue &);
    std::list<T> queue;
    std::mutex myMutex;
    // std::conditionVariable myCv;
};
...
```


Producer/Consumer

▪ `synchronized_queue.h` (2/2)

```
template<typename T>
void SynchronizedQueue<T>::put(const T & data)
{
    std::unique_lock<std::mutex> lck(myMutex);
    queue.push_back(data);
    //myCv.notify_one();
}

template<typename T>
T SynchronizedQueue<T>::get() {
    for(;;) {
        std::unique_lock<std::mutex> lck(myMutex);
        if(queue.empty()) continue;
        T result=queue.front();
        queue.pop_front();
        return result;
    }
}

#endif // SYNC_QUEUE_H_
```

Producer/Consumer

- What if we do not use the condition variable?
- The consumer is left “spinning” when the queue is empty
 - This takes up precious CPU cycles and slows down other threads in the system
 - Keeping the CPU busy increases power consumption
- Although the code is correct from a functional point of view this is a bad programming approach
 - When a thread has nothing to do it should block to free the CPU for other threads and reduce power consumption
- Extension: *Try to implement the version with a limited queue size*
 - The producer shall block if the queue reaches the maximum size

Active Object

- To instantiate “task objects”
- A thread function has no explicit way for other threads to communicate with it
 - Often data is passed to thread by global variables
- Conversely, this pattern allows us to wrap a thread into an object, thus having a “thread with methods you can call”
 - We may have member functions to pass data while the task is running, and collect results
- In some programming languages (e.g., Smaltalk and Objective C) all objects are “active objects”

Active Object

- The class includes a thread object and a member function **run()** implementing the task

```
#ifndef ACTIVE_OBJ_H_
#define ACTIVE_OBJ_H_
#include <atomic>
#include <thread>
class ActiveObject {
public:
    ActiveObject();
    ~ActiveObject();
private:
    virtual void run();
    ActiveObject(const ActiveObject &);
    ActiveObject& operator=(const ActiveObject &);
protected:
    std::thread t;
    std::atomic<bool> quit;
};
#endif // ACTIVE_OBJ_H_
```

Active Object

```
#include "active_object.h"
#include <chrono>
#include <functional>
#include <iostream>
using namespace std;
using namespace std::chrono;

ActiveObject::ActiveObject() :
    t(bind(&ActiveObject::run, this)), quit(false) {}

void ActiveObject::run() {
    while(!quit.load()) {
        cout<<"Hello world"<<endl;
        this_thread::sleep_for(milliseconds(500));
    }
}

ActiveObject::~~ActiveObject() {
    if(quit.load()) return; //For derived classes
    quit.store(true);
    t.join();
}
```

Active Object

- The constructor initialize the thread object, while the destructor takes care of joining it
- The **run ()** member function acts as a “main” concurrently executing
- Commonly used to implement threads communicating through a producer/consumer approach
- In the example implementation we used the “atomic” variable **quit** to terminate the **run ()** function when the object is destroyed
 - A normal boolean variable with a mutex would work as well

bind and function

- C has no way to decouple function *arguments binding* from the call
- In C++11 **bind** and **function** allow us to package a function and its arguments, and call it later

```
#include <iostream>
#include <functional>
using namespace std;


void printAdd(int a, int b) {
    cout<<a<<'+'<<b<<'='<<a+b<<endl;
}

int main() {
    function<void ()> func;
    func = bind(&printAdd, 2, 3);
    ...
    func();
}
```

We want to handle a function as if it was an object



We specify the function arguments without performing the call



Function call (with already packaged arguments)



Reactor

- The goal is to decouple the task creation from the execution
- A executor thread waits on a task queue
- Any other part of the program can push tasks into the queue
- Task are executed sequentially

The simplest solution is usually in a FIFO order

We are free to add to the “reactor” alternative thread scheduling functions

- C++11 **bind** and **function** allows us to create the task, leaving the starting time to the executor thread, in a second step

Reactor

- The class derives from `ActiveObject` to implement the executor thread and uses the `SynchronizedQueue` for the task queue

```
#ifndef REACTOR_H_
#define REACTOR_H_
#include <functional>
#include "synchronized_queue.h"
#include "active_object.h"

class Reactor: public ActiveObject {
public:
    void pushTask(std::function<void ()> func);
    virtual void ~Reactor();
private:
    virtual void run();
    SynchronizedQueue<std::function<void ()>> tasks;
};
#endif // REACTOR_H_
```

Reactor

```
#include "reactor.h"
using namespace std;

void doNothing() {}

void Reactor::pushTask(function<void ()> func) {
    tasks.put(func);
}

Reactor::~Reactor() {
    quit.store(true);
    pushTask(&doNothing);
    t.join(); // Thread derived from ActiveObject
}

void Reactor::run() {
    while(!quit.load())
        tasks.get()(); // Get a function and call it
}
```

Reactor

- In the example we are pushing a task to execute the `printAdd()` function

```
#include <iostream>

using namespace std;

void printAdd(int a, int b)
{
    cout<<a<<'+'<<b<<'='<<a+b<<endl;
}

int main()
{
    Reactor reac;
    reac.pushTask(bind(&printAdd, 2, 3));
    ...
}
```

ThreadPool

- The limit of the Reactor pattern is due to the fact that tasks are processed sequentially

The latency of the task execution is dependent on the length of the task queue

- To reduce latency and exploit multi-core computing architectures we can have multiple executor threads waiting on the same task queue
- *Try to implement your own ThreadPool at home!*